

lab2_2117902

October 16, 2024

1 Application of ML-based algorithm on The UNSW_NNB15 Datasets

1.1 Library

```
[1]: ##pip install scikit-learn pandas numpy prettyprint cupy tqdm matplotlib  
↳colorama
```

```
[2]: # Importing necessary libraries  
import pandas as pd  
import numpy as np  
from typing import Literal  
from sklearn.metrics import roc_curve, precision_recall_curve, auc  
from sklearn.preprocessing import LabelEncoder  
from sklearn.decomposition import PCA  
from sklearn.tree import DecisionTreeClassifier as DTC  
from sklearn.neighbors import KNeighborsClassifier  
from gc import collect  
from time import sleep  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import accuracy_score  
import cupy as cp  
from tqdm import tqdm  
import matplotlib.pyplot as plt  
import seaborn as sns  
from colorama import Style
```

1.2 Loading the data set

https://unsw-my.sharepoint.com/personal/z5025758_ad_unsw_edu_au/_layouts/15/onedrive.aspx?id=%2Fpersonal%2Fz5025758_ad_unsw_edu_au%2FDocuments%2FUnsw_NB15_training-set.csv

```
[3]: TRAIN_DATASET = 'UNSW_NB15_training-set.csv'  
TEST_DATASET = 'UNSW_NB15_testing-set.csv'
```

```
[4]: # Loading the csv data in a DataFrame  
df_train = pd.read_csv(TRAIN_DATASET)  
df_test = pd.read_csv(TEST_DATASET)
```

```
[5]: df_train.head(5)
```

```
[5]:
```

	id	dur	proto	service	state	spkts	dpkts	sbytes	dbytes	rate	\
0	1	0.121478	tcp	-	FIN	6	4	258	172	74.087490	
1	2	0.649902	tcp	-	FIN	14	38	734	42014	78.473372	
2	3	1.623129	tcp	-	FIN	8	16	364	13186	14.170161	
3	4	1.681642	tcp	ftp	FIN	12	12	628	770	13.677108	
4	5	0.449454	tcp	-	FIN	10	6	534	268	33.373826	

	...	ct_dst_sport_ltm	ct_dst_src_ltm	is_ftp_login	ct_ftp_cmd	\
0	...	1	1	0	0	
1	...	1	2	0	0	
2	...	1	3	0	0	
3	...	1	3	1	1	
4	...	1	40	0	0	

	ct_flw_http_mthd	ct_src_ltm	ct_srv_dst	is_sm_ips_ports	attack_cat	\
0	0	1	1	0	Normal	
1	0	1	6	0	Normal	
2	0	2	6	0	Normal	
3	0	2	1	0	Normal	
4	0	2	39	0	Normal	

	label
0	0
1	0
2	0
3	0
4	0

[5 rows x 45 columns]

```
[6]: def print_dataframe_shape(df:pd.DataFrame,name):print(f"The shape of {name} is:␣
↪{df.shape}")
```

```
[7]: print_dataframe_shape(df_train,'Training Set')
print_dataframe_shape(df_test,'Testing Set')
```

The shape of Training Set is: (175341, 45)

The shape of Testing Set is: (82332, 45)

```
[8]: # Creating a type of all the Features
Features = Literal['dur',
'proto',
'service',
'state',
'spkts',
```

```

'dpkts',
'sbytes',
'dbytes',
'rate',
'sttl',
'dttl',
'sload',
'dload',
'sloss',
'dloss',
'sinpkt',
'dinpkt',
'sjit',
'djit',
'swin',
'stcpb',
'dtcpb',
'dwin',
'tcprtt',
'synack',
'ackdat',
'smean',
'dmean',
'trans_depth',
'response_body_len',
'ct_srv_src',
'ct_state_ttl',
'ct_dst_ltm',
'ct_src_dport_ltm',
'ct_dst_sport_ltm',
'ct_dst_src_ltm',
'is_ftp_login',
'ct_ftp_cmd',
'ct_flw_http_mthd',
'ct_src_ltm',
'ct_srv_dst',
'is_sm_ips_ports',
'label']

```

1.2.1 Understanding the data

```

[9]: y_label = 'label' # setting the label
     features_list = df_train.columns.tolist()

```

```

[10]: def separate(df:pd.DataFrame,):
        features = features_list.copy()
        features.remove(y_label)

```

```

    return df.drop(y_label,axis=1),df[y_label] # separate a DataFrame of the
    ↪ features and the label

```

1.3 Preprocessing the Data

1.3.1 Cleaning the Data

Helper Function

```

[11]: def standardize(df,column:Features): # Standardizing the values of a features
    ↪ that has continuous values
        col_values = df[column].values

        mean = np.mean(col_values)
        std = np.std(col_values)
        col_values = col_values-mean
        col_values= col_values/std
        return pd.Series(col_values, name=column)

def min_max_scaling(df, column:Features): # MinMax scaling the values of a
    ↪ features that has continuous values
        col_values = df[column].values
        min_value = np.min(col_values)
        max_value = np.max(col_values)
        scaled_values = (col_values - min_value) / (max_value - min_value)
        return pd.Series(scaled_values, name=column)

def state_to_mask(state_vector: np.ndarray): # Creating a Mask for feature that
    ↪ has equally important state
        unique_val = np.unique(state_vector)
        size = len(unique_val)
        return { unique_val[mask]:mask for mask in range(size)}

def one_hot_encoding(state_mask:dict[int,str]): # Creating a function that
    ↪ returns a one hot encoding from a Mask
        def wrapper(mask: str):
            v = np.zeros((1, len(state_mask)))
            mask = state_mask[mask]
            v[0][mask] = 1
            return v
        return wrapper

def one_hot_vector_distance(v1: np.ndarray, v2: np.ndarray): # Compute the
    ↪ distance between two one hot encoding vector
        if v1.shape != v2.shape:
            raise
        if np.array_equal(v1, v2):

```

```

        return 0
    return 1

def str_encoder(df:pd.DataFrame,column:Features): # Set string to a equally
    ↪important state
    label_encoder = LabelEncoder()
    df[column] = label_encoder.fit_transform(df[column])

```

```

[12]: def remove_unecessaryFeature(df: pd.DataFrame, features: list = []): # Function
    ↪to remove features
    try:
        return df.drop(features, axis=1)
    except KeyError :
        return df

```

Cleaning Function ...

```

[13]: features_to_normalize=['dur','spkts','stcpb','dtpcb','dpkts','dbytes','sbytes',
                             'rate',

    ↪
    ↪'sload','dload','sloss','dloss','sinpkt','dinpkt','sjit','djit','tcprrt','synack','smean','
features_to_ohe=['proto','service','state','is_ftp_login','ct_ftp_cmd','ct_flw_http_mthd','ct_
initial_features_to_remove = ['id','attack_cat']

text_featuresType = ['proto','service','state']

def preprocess_final(df:pd.DataFrame, normalize:
    ↪Literal['min_max_scaling','standardize'],features_to_remove:
    ↪list=[]):

    ftr = set(features_to_remove)
    df
    ↪=remove_unecessaryFeature(df,[*initial_features_to_remove,*features_to_remove])
    ↪# Removing unecessary feature

    # Set values to a equally important state
    for feature in set(features_to_ohe).difference(ftr).
    ↪union(text_featuresType):
        str_encoder(df,feature)

    #Normalize continuous feature
    for feature in set(features_to_normalize).difference(ftr):
        df[feature] = normalize(df,feature)

    return df

```

```
[14]: def preprocess_partial(df:pd.DataFrame):

    # Removing unnecessary feature
    df =remove_unnecessaryFeature(df,['label',*initial_features_to_remove])

    # Set values to a equally important state
    for feature in text_featuresType:
        str_encoder(df,feature)

    #Normalize continuous feature
    for feature in features_to_normalize:
        df[feature] = standardize(df,feature)

    return df
```

1.4 Feature Selection

Looking for the features that has highest impact

```
[15]: df_feature_analysis= preprocess_partial(df_train) # Creating a preprocessed
↳DataFrame to compute some Feature Engineering
```

Correlation Matrix

```
[16]: # Find the highest correlation of a feature
def find_highest_correlation(corr_matrix:pd.DataFrame, target_feature:str):
    target_corr = corr_matrix[target_feature].drop(target_feature)
    highest_corr_feature = target_corr.idxmax()
    highest_corr_value = target_corr[highest_corr_feature]

    return highest_corr_feature, highest_corr_value
```

```
[17]: corr_matrix = df_feature_analysis.corr().apply(abs) # Get the correlation
↳matrix with absolute values
corr_matrix
```

```
[17]:
```

	dur	proto	service	state	spkts	dpkts	\
dur	1.000000	0.124502	0.008234	0.103443	0.254559	0.181182	
proto	0.124502	1.000000	0.170032	0.172441	0.013469	0.026439	
service	0.008234	0.170032	1.000000	0.144978	0.114403	0.077338	
state	0.103443	0.172441	0.144978	1.000000	0.078701	0.098268	
spkts	0.254559	0.013469	0.114403	0.078701	1.000000	0.390067	
dpkts	0.181182	0.026439	0.077338	0.098268	0.390067	1.000000	
sbytes	0.199731	0.005920	0.105188	0.049300	0.963791	0.188476	
dbytes	0.144134	0.015812	0.035492	0.059759	0.206609	0.971907	
rate	0.120966	0.013924	0.141709	0.432307	0.076358	0.098202	
sttl	0.012196	0.049944	0.295302	0.584697	0.102723	0.192580	
dttl	0.044159	0.113184	0.262970	0.375533	0.068246	0.053861	

sload	0.081749	0.004759	0.166339	0.292570	0.051646	0.066710
dload	0.050603	0.046375	0.099581	0.150501	0.075897	0.139145
sloss	0.198597	0.011392	0.114522	0.060125	0.971069	0.204883
dloss	0.142963	0.020002	0.051495	0.071056	0.207798	0.978636
sinpkt	0.080055	0.562789	0.089971	0.095492	0.017587	0.022160
dinpkt	0.152142	0.052417	0.020190	0.076235	0.001678	0.006514
sjit	0.144413	0.016011	0.011469	0.045441	0.000384	0.000229
djit	0.157443	0.019388	0.090262	0.064747	0.017096	0.054371
swin	0.022047	0.138967	0.292887	0.367493	0.131813	0.183703
stcpb	0.013183	0.108571	0.237103	0.314361	0.107410	0.144119
dtcpb	0.014724	0.108630	0.237723	0.313922	0.102161	0.142667
dwin	0.017527	0.137605	0.300035	0.397710	0.133102	0.185555
tcprrt	0.053125	0.079193	0.140239	0.278469	0.039187	0.020915
synack	0.051093	0.073528	0.110995	0.261882	0.035507	0.015936
ackdat	0.049332	0.076362	0.155811	0.264946	0.038725	0.023899
smean	0.090028	0.042157	0.224861	0.070796	0.216592	0.014697
dmean	0.025336	0.077296	0.145641	0.256392	0.150237	0.441445
trans_depth	0.002071	0.020709	0.191839	0.056128	0.008834	0.029042
response_body_len	0.078915	0.006005	0.056951	0.025541	0.087217	0.442194
ct_srv_src	0.113709	0.203057	0.058269	0.385515	0.069127	0.079095
ct_state_ttl	0.186293	0.162433	0.205943	0.759825	0.086170	0.150023
ct_dst_ltm	0.086300	0.191101	0.047685	0.328748	0.060194	0.071909
ct_src_dport_ltm	0.094091	0.174965	0.038347	0.372309	0.068373	0.086695
ct_dst_sport_ltm	0.093923	0.165796	0.051106	0.408662	0.072484	0.094267
ct_dst_src_ltm	0.101760	0.175708	0.006774	0.429906	0.077553	0.094085
is_ftp_login	0.020641	0.018003	0.071051	0.051970	0.009951	0.013491
ct_ftp_cmd	0.020641	0.018003	0.071051	0.051970	0.009951	0.013491
ct_flw_http_mthd	0.024743	0.028809	0.266206	0.078856	0.006084	0.047974
ct_src_ltm	0.080871	0.168121	0.028599	0.323019	0.061584	0.075190
ct_srv_dst	0.115336	0.198594	0.048011	0.387446	0.069598	0.078342
is_sm_ips_ports	0.035370	0.585941	0.088847	0.094198	0.017770	0.021765

	sbytes	dbytes	rate	sttl	...	ct_dst_ltm \
dur	0.199731	0.144134	0.120966	0.012196	...	0.086300
proto	0.005920	0.015812	0.013924	0.049944	...	0.191101
service	0.105188	0.035492	0.141709	0.295302	...	0.047685
state	0.049300	0.059759	0.432307	0.584697	...	0.328748
spkts	0.963791	0.206609	0.076358	0.102723	...	0.060194
dpkts	0.188476	0.971907	0.098202	0.192580	...	0.071909
sbytes	1.000000	0.009926	0.028468	0.020860	...	0.026661
dbytes	0.009926	1.000000	0.059475	0.135515	...	0.042633
rate	0.028468	0.059475	1.000000	0.407572	...	0.317229
sttl	0.020860	0.135515	0.407572	1.000000	...	0.271383
dttl	0.063009	0.023559	0.414546	0.032823	...	0.381678
sload	0.018322	0.040430	0.602492	0.276475	...	0.076471
dload	0.007829	0.104757	0.153051	0.397431	...	0.100953
sloss	0.996109	0.017366	0.042923	0.044667	...	0.036965

dloss	0.006804	0.996504	0.075259	0.162628	...	0.054538
sinpkt	0.006565	0.013618	0.075745	0.206571	...	0.072241
dinpkt	0.000024	0.007701	0.051539	0.003215	...	0.042781
sjit	0.002054	0.002422	0.063370	0.022676	...	0.046592
djit	0.003516	0.047354	0.085802	0.123435	...	0.057296
swin	0.050450	0.113148	0.515681	0.416843	...	0.412379
stcpb	0.043164	0.086894	0.408750	0.337305	...	0.326216
dtcpb	0.037988	0.086453	0.409046	0.334114	...	0.327530
dwin	0.050981	0.114269	0.518117	0.424320	...	0.415255
tcprrt	0.043624	0.003907	0.300794	0.039777	...	0.286773
synack	0.039739	0.000101	0.279271	0.042590	...	0.264577
ackdat	0.042883	0.007546	0.290051	0.032293	...	0.278326
smean	0.232348	0.036635	0.113232	0.010029	...	0.162651
dmean	0.004973	0.419965	0.273323	0.550389	...	0.203729
trans_depth	0.003428	0.030912	0.078556	0.063904	...	0.069216
response_body_len	0.001620	0.470905	0.022752	0.050454	...	0.016102
ct_srv_src	0.034395	0.045529	0.357704	0.346079	...	0.841280
ct_state_ttl	0.012053	0.089944	0.431534	0.672325	...	0.302420
ct_dst_ltm	0.026661	0.042633	0.317229	0.271383	...	1.000000
ct_src_dport_ltm	0.026490	0.052135	0.353589	0.344104	...	0.962052
ct_dst_sport_ltm	0.027281	0.056901	0.390721	0.379930	...	0.870644
ct_dst_src_ltm	0.032061	0.054633	0.383094	0.404346	...	0.852252
is_ftp_login	0.004515	0.010460	0.068140	0.124157	...	0.048527
ct_ftp_cmd	0.004515	0.010460	0.068140	0.124157	...	0.048527
ct_flw_http_mthd	0.002185	0.051403	0.109297	0.112833	...	0.085540
ct_src_ltm	0.027479	0.045594	0.310876	0.273252	...	0.886072
ct_srv_dst	0.034553	0.044531	0.362883	0.340678	...	0.852583
is_sm_ips_ports	0.006367	0.013147	0.072948	0.220429	...	0.069455

	ct_src_dport_ltm	ct_dst_sport_ltm	ct_dst_src_ltm \
dur	0.094091	0.093923	0.101760
proto	0.174965	0.165796	0.175708
service	0.038347	0.051106	0.006774
state	0.372309	0.408662	0.429906
spkts	0.068373	0.072484	0.077553
dpkts	0.086695	0.094267	0.094085
sbytes	0.026490	0.027281	0.032061
dbytes	0.052135	0.056901	0.054633
rate	0.353589	0.390721	0.383094
sttl	0.344104	0.379930	0.404346
dttl	0.366308	0.389429	0.403465
sload	0.100118	0.082462	0.155030
dload	0.143573	0.153429	0.161192
sloss	0.039158	0.041109	0.045857
dloss	0.066411	0.072203	0.070905
sinpkt	0.060851	0.057659	0.081595
dinpkt	0.038731	0.039644	0.042856

sjit	0.043927	0.045747	0.047338
djit	0.071165	0.075841	0.081518
swin	0.453084	0.497973	0.492118
stcpb	0.353927	0.389607	0.394579
dtcpb	0.354069	0.389581	0.394566
dwin	0.448574	0.493572	0.499716
tcprrt	0.264803	0.280883	0.290241
synack	0.244511	0.260560	0.265808
ackdat	0.256785	0.271092	0.283801
smean	0.169091	0.199021	0.152219
dmean	0.244463	0.264171	0.279326
trans_depth	0.069969	0.073894	0.084412
response_body_len	0.020611	0.021715	0.021633
ct_srv_src	0.866010	0.823583	0.967138
ct_state_ttl	0.353778	0.393091	0.427705
ct_dst_ltm	0.962052	0.870644	0.852252
ct_src_dport_ltm	1.000000	0.906793	0.869941
ct_dst_sport_ltm	0.906793	1.000000	0.838678
ct_dst_src_ltm	0.869941	0.838678	1.000000
is_ftp_login	0.064055	0.065305	0.062574
ct_ftp_cmd	0.064055	0.065305	0.062574
ct_flw_http_mthd	0.085699	0.086594	0.106129
ct_src_ltm	0.897438	0.803013	0.783753
ct_srv_dst	0.868850	0.830152	0.972370
is_sm_ips_ports	0.056858	0.053224	0.079765

	is_ftp_login	ct_ftp_cmd	ct_flw_http_mthd	ct_src_ltm	\
dur	0.020641	0.020641	0.024743	0.080871	
proto	0.018003	0.018003	0.028809	0.168121	
service	0.071051	0.071051	0.266206	0.028599	
state	0.051970	0.051970	0.078856	0.323019	
spkts	0.009951	0.009951	0.006084	0.061584	
dpkts	0.013491	0.013491	0.047974	0.075190	
sbytes	0.004515	0.004515	0.002185	0.027479	
dbytes	0.010460	0.010460	0.051403	0.045594	
rate	0.068140	0.068140	0.109297	0.310876	
sttl	0.124157	0.124157	0.112833	0.273252	
dttl	0.107208	0.107208	0.223652	0.365404	
sload	0.046194	0.046194	0.073920	0.084412	
dload	0.027810	0.027810	0.039246	0.098149	
sloss	0.005688	0.005688	0.002049	0.038795	
dloss	0.007763	0.007763	0.048869	0.057412	
sinpkt	0.014458	0.014458	0.018829	0.081130	
dinpkt	0.002255	0.002255	0.046655	0.042445	
sjit	0.005798	0.005798	0.088052	0.045108	
djit	0.081014	0.081014	0.100563	0.062372	
swin	0.129555	0.129555	0.207313	0.402189	

stcpb	0.097536	0.097536	0.161696	0.318968
dtcpb	0.100410	0.100410	0.172032	0.317651
dwin	0.130834	0.130834	0.209360	0.403987
tcprrt	0.067715	0.067715	0.163332	0.278013
synack	0.056794	0.056794	0.144906	0.256475
ackdat	0.071807	0.071807	0.164719	0.269845
smean	0.043298	0.043298	0.017901	0.159718
dmean	0.023999	0.023999	0.129436	0.201604
trans_depth	0.016177	0.016177	0.226152	0.065314
response_body_len	0.004691	0.004691	0.065238	0.018091
ct_srv_src	0.089827	0.089827	0.120111	0.781051
ct_state_ttl	0.075058	0.075058	0.097270	0.296638
ct_dst_ltm	0.048527	0.048527	0.085540	0.886072
ct_src_dport_ltm	0.064055	0.064055	0.085699	0.897438
ct_dst_sport_ltm	0.065305	0.065305	0.086594	0.803013
ct_dst_src_ltm	0.062574	0.062574	0.106129	0.783753
is_ftp_login	1.000000	1.000000	0.022505	0.046326
ct_ftp_cmd	1.000000	1.000000	0.022505	0.046326
ct_flw_http_mthd	0.022505	0.022505	1.000000	0.074768
ct_src_ltm	0.046326	0.046326	0.074768	1.000000
ct_srv_dst	0.087511	0.087511	0.118709	0.777891
is_sm_ips_ports	0.015003	0.015003	0.024007	0.078886

	ct_srv_dst	is_sm_ips_ports
dur	0.115336	0.035370
proto	0.198594	0.585941
service	0.048011	0.088847
state	0.387446	0.094198
spkts	0.069598	0.017770
dpkts	0.078342	0.021765
sbytes	0.034553	0.006367
dbytes	0.044531	0.013147
rate	0.362883	0.072948
sttl	0.340678	0.220429
dttl	0.431188	0.091137
sload	0.141168	0.049327
dload	0.087247	0.035069
sloss	0.045459	0.009492
dloss	0.058605	0.016669
sinpkt	0.086988	0.941319
dinpkt	0.045648	0.011306
sjit	0.049711	0.013987
djit	0.082087	0.018827
swin	0.466245	0.115622
stcpb	0.373117	0.090374
dtcpb	0.374050	0.090525
dwin	0.473821	0.114671

tcprrtt	0.316713	0.065994
synack	0.289605	0.061274
ackdat	0.310164	0.063635
smean	0.171203	0.056094
dmean	0.227964	0.060813
trans_depth	0.093901	0.017258
response_body_len	0.027303	0.005004
ct_srv_src	0.980323	0.088456
ct_state_ttl	0.364200	0.092616
ct_dst_ltm	0.852583	0.069455
ct_src_dport_ltm	0.868850	0.056858
ct_dst_sport_ltm	0.830152	0.053224
ct_dst_src_ltm	0.972370	0.079765
is_ftp_login	0.087511	0.015003
ct_ftp_cmd	0.087511	0.015003
ct_flw_http_mthd	0.118709	0.024007
ct_src_ltm	0.777891	0.078886
ct_srv_dst	1.000000	0.085149
is_sm_ips_ports	0.085149	1.000000

[42 rows x 42 columns]

```
[18]: corr_feature={}
      corr_tresh= 0.94

      for feature in df_feature_analysis.columns:
          #Get the highest correlation for each feature
          f,score= find_highest_correlation(corr_matrix,feature)
          #Loading the feature and the score
          corr_feature[feature] = {'feature':f, 'score':score}

      #Removing duplicates correlation since the corr matrix is a triangular matrix
      for feature in df_feature_analysis.columns:
          try:
              t = corr_feature[feature]['feature']
              if corr_feature[t]['feature'] == feature:
                  del corr_feature[t]
          except KeyError:
              continue

      # Sorting by the highest correlation value
      corr_feature = dict(sorted(corr_feature.items(), key=lambda item:
          ↪item[1]['score'], reverse=True))

      # Filtering all correlation that is not above the threshold value
      corr_feature = {f:corr_feature[f] for f in corr_feature.keys() if
          ↪corr_feature[f]['score'] >= corr_tresh }
      corr_feature
```

```
[18]: {'is_ftp_login': {'feature': 'ct_ftp_cmd', 'score': 1.0},
      'dbytes': {'feature': 'dloss', 'score': 0.996503594762374},
      'sbytes': {'feature': 'sloss', 'score': 0.9961094729147967},
      'swin': {'feature': 'dwin', 'score': 0.9901399299415929},
      'ct_srv_src': {'feature': 'ct_srv_dst', 'score': 0.9803230099911133},
      'dpkts': {'feature': 'dloss', 'score': 0.9786363765710283},
      'ct_dst_src_ltm': {'feature': 'ct_srv_dst', 'score': 0.9723704538697349},
      'spkts': {'feature': 'sloss', 'score': 0.9710686917738162},
      'ct_dst_ltm': {'feature': 'ct_src_dport_ltm', 'score': 0.9620518416459877},
      'tcprrt': {'feature': 'synack', 'score': 0.9494676611067793},
      'ackdat': {'feature': 'tcprrt', 'score': 0.941760373812716},
      'sinpkt': {'feature': 'is_sm_ips_ports', 'score': 0.941318900735516}}
```

```
[19]: def compare_features(feature:list[Features],order_by:Features, ascending:
      ↪bool=True):
      if y_label not in feature:
          feature.append(y_label)
          print(feature)
      if order_by not in feature:
          raise ValueError('order_by must be in the feature parameter')

      if order_by == y_label:
          raise ValueError(f'cannot order_by the {y_label}')
      return df_train[feature].sort_values(by=order_by,axis=0,ascending=ascending)
```

`is_ftp_login` and `ct_ftp_cmd` are exactly the same, so we can remove one them. For the rest we determined by looking directly at the `df_train` data and concludes that they are indeed highly correlated and prediction resulting removing will not affect that much

```
[20]: features_to_remove=['is_ftp_login']
      #features_to_remove=['is_ftp_login','sbytes','dbytes','swin','dpkts','spkts']
```

PCA

```
[21]: top_n_components = 30
```

```
[22]: df_pca = df_feature_analysis.drop(features_to_remove,axis=1) #removing the
      ↪feature we did not need from the correlation analysis
      feature_cov = np.dot(df_pca.transpose(), df_pca)/len(df_pca)
      eigenvalues, eigenvectors = np.linalg.eig(feature_cov) # Getting the eigenvalues
      pca_index= np.argsort(eigenvalues)[::-1][:top_n_components]
      pca_feature = df_feature_analysis.columns[pca_index]
      pca_feature # Getting the most important features (principal components)
      ↪ordered by the highest eigenvalues
```

```
[22]: Index(['dur', 'proto', 'service', 'state', 'spkts', 'dpkts', 'sbytes',
      'dbytes', 'rate', 'sttl', 'dttl', 'sload', 'dload', 'sloss', 'dloss',
      'sinpkt', 'dinpkt', 'sjit', 'djit', 'swin', 'stcpb', 'dtcpb', 'dwin',
```

```

        'tcprrtt', 'synack', 'ackdat', 'smean', 'dmean', 'trans_depth',
        'response_body_len'],
        dtype='object')

```

```

[23]: # Projecting the training dataset to the top_n_components SPACE to reduce the
      ↪ features and keep the most information and variance
def toPCA_space(df:pd.DataFrame,pca_list,top_n_components=top_n_components):
    pca = PCA(n_components=top_n_components)
    pca.fit(df.values)
    pca_data = pca.transform(df.values)
    return pd.DataFrame(pca_data, columns=pca_list)

```

1.4.1 Final Preprocessing Step

Based on the various technique we decided to remove those features from the features that has equally important state

```

[24]: features_to_ohe = list(set(features_to_ohe).difference(features_to_remove))
      ↪ #features_to_ohe = []
      features_to_ohe

```

```

[24]: ['is_sm_ips_ports',
      'state',
      'ct_state_ttl',
      'service',
      'ct_flw_http_mthd',
      'proto',
      'ct_ftp_cmd']

```

```

[ ]: # Preprocessing the dataset and separate them to do the a Model training
X_train, Y_train =
    ↪ separate(preprocess_final(df_train,features_to_remove=features_to_remove))
X_test, Y_test =
    ↪ separate(preprocess_final(df_test,features_to_remove=features_to_remove))
X_train_PCA = toPCA_space(X_train,pca_feature.to_list(),)
X_test_PCA = toPCA_space(X_test,pca_feature.to_list(),)

```

```

[26]: # NOTE uncommenting to put the dataset in the PCA space
      ↪ #X_train = X_train_PCA
      ↪ #X_test = X_test_PCA
      print_dataframe_shape(X_train,'Training Set')
      print_dataframe_shape(X_test,'Testing Set')

```

The shape of Training Set is: (175341, 41)

The shape of Testing Set is: (82332, 41)

```
[27]: #After this step those values are now unimportant
del df_train, df_test, df_feature_analysis, corr_matrix
collect()
```

[27]: 20

1.5 Model

```
[28]: # Creating a Label class that store the metadata of binary classification
class LabelMetadata:

    def __init__(self, pos_class, ␣
↪neg_class, pos_name, neg_name, preferred_class=None) -> None:
        self.PositiveClass:int = pos_class
        self.NegativeClass:int = neg_class
        self.NegativeName:str = neg_name
        self.PositiveName:str = pos_name
        self.PreferredClass= self.NegativeClass if preferred_class is None else ␣
↪self.PositiveClass

        self.answer={
            self.PositiveClass: self.PositiveName,
            self.NegativeClass: self.NegativeName
        }

problem_label_class = LabelMetadata(0,1,'Normal','Attack')
```

```
[64]: # Base class of the Binary Classifier
class BinaryClassifier:

    # setting up the variables in the init function
    def __init__(self, label_class=problem_label_class):
        self.X = None
        self.Y = None
        self.Y_Pred:list = None
        self.Y_PredProba=[]
        self.label_class = label_class

    def fit(self):
        ...

    def predict(self):
        ...

    # Method to count the label class in a vector
    def _label_count(self, label_vectors):
        n = len(label_vectors)
```

```

sum_one = list(label_vectors).count(1)
sum_zero = n - sum_one
return sum_zero, sum_one, n

# Given the Y_test, calculate the True Positive, True Negative, False_
↪ Positive, False Negative and other info
def _compute_analysis(self, y_test):
    self.TP = 0
    self.TN = 0
    self.FP = 0
    self.FN = 0

    for truth, pred in zip(y_test, self.Y_Pred):
        if truth == self.label_class.NegativeClass and pred == self.
↪ label_class.NegativeClass:
            self.TN += 1
        elif truth == self.label_class.PositiveClass and pred == self.
↪ label_class.PositiveClass:
            self.TP += 1
        elif truth == self.label_class.NegativeClass and pred == self.
↪ label_class.PositiveClass:
            self.FP += 1
        else:
            self.FN += 1
    self.roc_info = roc_curve(y_test, self.Y_PredProba)
    self.precision_recall_info = precision_recall_curve(y_test, self.
↪ Y_PredProba)

def plot_confusion_matrix(self):
    confusion_matrix = np.array([[self.TN, self.FP], [self.FN, self.TP]])

    plt.figure(figsize=(6, 4))
    sns.heatmap(confusion_matrix, annot=True, fmt='d', cmap='Blues',
↪ cbar=False,
                xticklabels=['Predicted Negative', 'Predicted Positive'],
                yticklabels=['Actual Negative', 'Actual Positive'])

    plt.title('Confusion Matrix')
    plt.xlabel('Predicted Labels')
    plt.ylabel('True Labels')
    plt.show()

def plot_roc_curve(self):
    fpr, tpr, thresholds_roc = self.roc_info
    plt.figure()

```

```

plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area =_{
↪{auc(fpr, tpr):0.2f}})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc="lower right")
plt.show()

def plot_precision_recall_curve(self):
    precision, recall, thresholds_pr = self.precision_recall_info
    plt.figure()
    plt.plot(recall, precision, color='b', lw=2)
    plt.xlabel('Recall')
    plt.ylabel('Precision')
    plt.title('Precision-Recall Curve')
    plt.show()

@property
def accuracy(self):
    return (self.TP + self.TN)/(self.TP + self.TN +self.FP + self.FN)

@property
def f_score(self):
    return (2*self.precision * self.recall)/(self.precision+self.recall)

@property
def precision(self):
    return (self.TP)/(self.TP + self.FP)

@property
def recall(self):
    return self.TP/(self.TP + self.FN)

```

1.5.1 Decision Tree

Question

```

[30]: # Encapsulate the concept of the best split into a Question object, since the_
↪best split is a question of which value will determine the split
class Question:

    def __init__(self,feature:Features,value:float,information_gain:float):
        self.feature = feature
        self.value = value
        self.information_gain = information_gain

    def split(self,dataset:pd.DataFrame)->tuple[pd.DataFrame,pd.DataFrame]:

```



```

...

#Representing the object in a Question form
def _repr(self, _type):
    return f'Is {Style.DIM}{self.feature}{Style.RESET_ALL} {Style.
↪BRIGHT}{_type}{Style.RESET_ALL} to {Style.DIM}{self.value:0.5f}{Style.
↪RESET_ALL} ? - Gain[{self.information_gain:.6f}]'

def match(self, vector:pd.Series) -> bool:
    ...

def __str__(self):
    return self.__repr__()

# Dunder method to help the programming

def __eq__(self, other):
    return self.information_gain == other.information_gain

def __ne__(self, other):
    return self.information_gain != other.information_gain

def __gt__(self, other):
    return self.information_gain > other.information_gain

def __ge__(self, other):
    return self.information_gain >= other.information_gain

# When the feature has equally important state
class QuestionEqual(Question):

    # Splitting a dataset by asking if the feature is equal or not equal to the
    ↪value
    def split(self, dataset:pd.DataFrame):
        return dataset[dataset[self.feature] == self.value], ↪
    ↪dataset[dataset[self.feature] != self.value]

    # Whether the feature of a single vector answer the questions
    def match(self, vector):
        return vector[self.feature] == self.value

    def __repr__(self):
        return super().__repr__('equal')

class QuestionThresh(Question):

```

```

    # Splitting a dataset by asking if the feature is greater equal or
    ↪ strictly lower than the threshold(value)
    def split(self, dataset:pd.DataFrame):
        return dataset[dataset[self.feature] >= self.value],
    ↪ dataset[dataset[self.feature] < self.value]

    # Whether the feature of a single vector answer the questions
    def match(self,vector):
        return vector[self.feature] >= self.value

    def __repr__(self):
        return super().__repr__('greater or equal')

```

Node class

```

[31]: # Base Node class
class Node:
    node_count=0
    def __init__(self):
        Node.node_count+=1

# Node that has other node basically a subtree
class TreeNode(Node):
    def __init__(self,question:Question,left:Node,right:Node):
        super().__init__()
        self.question = question
        """
        Satisfy the match
        """
        self.left=left
        """
        Dissatisfy the match
        """
        self.right = right

    # Ask the question of a certain vector and dirige to the child that has the
    ↪ answer
    def match(self,vector) -> Node:
        return self.left if self.question.match(vector) else self.right

    def __repr__(self):
        return repr(self.question)

# Label that has the final answer of a series of question
class LeafNode(Node):

    def __init__(self,probabilities,label_class:LabelMetadata):

```

```

    super().__init__()
    self.proba = probabilities
    self.label_class = label_class
    self.answer:Literal[0,1,None] = self._compute_answer()

    # Compute an answer based on a binary probabilities
    def _compute_answer(self):
        label_0 = self.proba[0]
        label_1 = self.proba[1]

        if label_0 == label_1:
            return self.label_class.PreferredClass

        return 1 if label_1 > label_0 else 0

    @property
    def answer_proba(self):
        return self.proba[self.answer] # probability of a final answer

    def __repr__(self):
        return f'It is {Style.DIM}{self.label_class.answer[self.answer]}{Style.
↪RESET_ALL} with a probability of {self.answer_proba:.4f} %'

```

DecisionTree class

```

[32]: x_train, x_val, y_train, y_val = train_test_split(X_train, Y_train, test_size=0.
↪2, random_state=42)

```

```

[63]: ImpurityType=Literal['gini_index','entropy']

class DecisionTreeClassifier(BinaryClassifier):

    # Constructor of that set all the hyperparameters
    def __init__(self,max_height:int,min_information_gain:float,min_sample:
↪int,impurity:
↪ImpurityType='entropy',label_class=problem_label_class,ohe_feature=features_to_ohe):
↪
        super().__init__(label_class)
        self.max_height = max_height if max_height is not None else float('inf')
        self.min_information_gain = min_information_gain
        self.min_sample = min_sample
        self.impurity_type = impurity
        self.ohe_features = ohe_feature
        self.impurity = self._gini_impurity if impurity == 'gini_index' else self.
↪_entropy
        self.root: TreeNode = None

```

```

# loading the dataset to fit the model
def fit(self,x_train:pd.DataFrame,y_train:pd.DataFrame):
    self.X = x_train
    self.Y = y_train
    self.root_dataset = pd.concat([self.X,self.Y],axis=1)

    # Train a model using a training set and predict the values of a validation
    ↪set if they are specified
    def train(self,x_val=None,y_val=None):

        self.root =self._build_tree(self.root_dataset)
        if x_val is None or y_val is None:
            return
        self.predict(x_val,y_val)

    # Return the representation of the Decision Tree Classifier
    def __str__(self) -> str:
        return self.__repr__()

    def __repr__(self):
        return f'DecisionTree(Max_Depth={self.max_height},Min_Inf_Gain={self.
    ↪min_information_gain},Min_Sample={self.min_sample},Impurity={self.
    ↪impurity_type})'

    # Predict values of a testing/validation test and compute analysis to aid
    ↪several metrics
    def predict(self,x_to_test,y_to_test):
        self.Y_Pred = self._predict(x_to_test)
        self._compute_analysis(y_to_test)

    # from a dataframe, each vector will traverse the trained tree and have a
    ↪computed answer, then load the self.Y_pred
    def _predict(self,x_to_test:pd.DataFrame):
        return x_to_test.apply(self._traverse_tree,axis=1).values

    # computes the entropy of a dataset
    def _entropy(self,labels:np.ndarray):
        try:
            probabilities = np.array(self._compute_target_probabilities(labels))
        except ZeroDivisionError:
            return float('inf')
        return -np.sum(probabilities * np.log2(probabilities))

    # compute the gini impurity of a dataset
    def _gini_impurity(self,labels:np.ndarray):
        try:

```

```

        probabilities = np.array(self._compute_target_probabilities(labels))
    except ZeroDivisionError:
        return float('inf')
    return (1 - np.sum(probabilities**2))

# get the probabilities of the binary class
def _compute_target_probabilities(self, labels_vectors):

    sum_zero, sum_one, n = self._label_count(labels_vectors)
    return [sum_zero/n, sum_one/n]

# compute the information gain
def _information_gain(self, current_information_gain: float, mean_impurity:
    float):
    return current_information_gain - mean_impurity

# build a tree recursively from a dataset, return a Leaf node if the a
condition match any hyperparameter otherwise go further into build another
TreeNode
def _build_tree(self, dataset: pd.DataFrame, current_depth: int = 0) -> TreeNode |
    LeafNode:

    parent_gain = self.impurity(dataset['label'].values) # current gain of
    the dataset
    current_n = len(dataset) # size of the current dataset

    if current_depth >= self.max_height or current_n < self.min_sample or
    parent_gain < self.min_information_gain: # if the dataset has any of this
    condition true then return a Leaf Node with an answer
        try:
            proba = self._compute_target_probabilities(dataset['label'].values)
        # get the probability of each class
        except ZeroDivisionError: # if theres a zero division error, return a
        node with the preferred class
            t = [0, 0]
            t[self.label_class.NegativeClass] = 1
            t[self.label_class.PositiveClass] = 0
            proba = np.array(t)
            return LeafNode(proba, self.label_class)

    best_question = self._find_best_split(dataset, parent_gain) # Get the best
    split by finding the best question

    left_dataset, right_dataset = self._split_dataset(dataset, best_question)
    # split the dataset into matching or no the question

```

```

        left_child = self._build_tree(left_dataset,current_depth+1) # if
↳matching build another tree
        right_child = self._build_tree(right_dataset,current_depth+1) # if not
↳matching build another tree

        return TreeNode(best_question,left_child,right_child)

    def _traverse_tree(self,x_vector:pd.Series):
        current_node:LeafNode | TreeNode = self.root
        while isinstance(current_node,TreeNode): # Asking questions(TreeNode)
↳till having an answer(LeafNode)
            current_node = current_node.match(x_vector)

        self.Y_PredProba.append(current_node.answer_proba) # loading the
↳probability of an answer
        return current_node.answer # Final answer

    def print_tree(self,): # printing the tree
        self._print_tree(self.root,0)

    def _print_tree(self,node:Node| TreeNode, depth,answer =None):
        print(' '*depth,' ' if answer is None else answer,node) # print the
↳question
        if type(node) is TreeNode:
            self._print_tree(node.left,depth+1,'YES...') # calling recursively
↳to the left
            self._print_tree(node.right,depth+1,'NO...') # calling recursively
↳to the right

    def _split_dataset(self,dataset:pd.DataFrame,question:Question): # splitting
↳the dataset into two based on the current best question
        return question.split(dataset)

    def _find_best_split(self,dataset:pd.DataFrame,current_gain:
↳float)->Question: # Fining the best question

        best_question = None
        for feature in dataset.columns:
            if feature == y_label: # the dataset contain the y_label value so we
↳skip it
                continue
            if feature in self.ohe_features:
                for values in dataset[feature].unique(): # going trough every unique
↳value if the state are equally important

```

```

        best_question = self._compute_best_question(dataset,
↪current_gain,feature, values,best_question,QuestionEqual) # compute the
↪global best question
        else: # if the feature has continuous values
            val_unique_mean = dataset[feature].unique().mean() # splitting by
↪mean from the unique values
            val_mean = dataset[feature].mean() # splitting by the mean
            val_median = dataset[feature].median() # splitting by the median

            for values in [val_unique_mean,val_mean,val_median]:
                best_question = self._compute_best_question(dataset,
↪current_gain,feature, values,best_question,QuestionThresh) # compute the
↪global best question

    return best_question

    def _compute_best_question(self, dataset:pd.DataFrame, current_gain:
↪float,feature:Features, values:float,best_question:Question,Q_type:type)
↪->Question:

        N = len(dataset)

        # splitting
        if Q_type== QuestionEqual: # if its a one hot encoding feature(ohe/
↪equally important state)
            y_satisfaction,y_dissatisfaction = dataset[dataset[feature]==
↪values].label.values,dataset[dataset[feature]!= values].label.values

            else: # if the feature is a continuous
                y_satisfaction, y_dissatisfaction = dataset[dataset[feature] >=
↪values].label.values, dataset[dataset[feature] < values].label.values

            #get the mean impurity from the previous split
            mean_impurity = (len(y_dissatisfaction)/N)*self.
↪impurity(y_dissatisfaction) + (len(y_satisfaction)/N)*self.
↪impurity(y_satisfaction)

            #calculating the information gain
            info_gain = self._information_gain(current_gain,mean_impurity)
            question = Q_type(feature,values,info_gain) # creating a question

            if best_question is None:
                return question # return the question if none best question were given

            return question if question > best_question else best_question # return
↪the question that maximize the information gain

```

```

# comparing decision tree model by the accuracy
def __eq__(self, other):
    return self.accuracy == other.accuracy

def __ne__(self, other):
    return self.accuracy != other.accuracy

def __gt__(self, other):
    return self.accuracy > other.accuracy

def __ge__(self, other):
    return self.accuracy >= other.accuracy

```

1.5.2 K-Nearest Neighbors

```

[62]: class KNNClassifier(BinaryClassifier):

    # Give the parameters aiding our KNN
    def __init__(self, ohe_feature: list[str], max_k: int = None, N_batch = 100) -> None:
        super().__init__()
        self.K = max_k
        self.N_batch = N_batch
        self.ohe_features = ohe_feature
        self.ohe_func = cp.vectorize(self._to_one_hot_encoding) # creating a vectorize function

    # loading the dataset and set a K value
    def fit(self, X_train: pd.DataFrame, Y_train: pd.DataFrame):
        self.X = X_train
        self.Y = Y_train
        self.x_num, self.x_ohe = self._split(X_train) # split the numerical values and ohe features
        if self.K is not None:
            self.K = self._to_odd_number(self.K - 1) # ensure that the K is an odd number
        else:
            self.K = self._to_odd_number(round(len(self.X) ** 0.5)) # give square root of len of the training dataset for a K (by convention)

    def _to_odd_number(self, val):
        return val - 1 if val % 2 == 0 else val # return an odd number if its even

    def _split(self, df: pd.DataFrame):
        return df.drop(self.ohe_features, axis=1), df[self.ohe_features] # split the numerical values and ohe features

```



```

def predict(self,x_test,y_test):
    dataframes_indices = self._predict(self._split(x_test)) # get the all K
    ↪closest indices
    self.df_distances = pd.DataFrame(pd.concat(dataframes_indices).apply(self.
    ↪_prevote,axis=1)) # concatenate and transform the indices into its specified
    ↪label
    self.df_distances.columns = ['label']
    self.Y_Pred = self.df_distances.label.apply(self._vote_majority).values
    ↪# vote the label
    del dataframes_indices, self.df_distances
    collect()
    self._compute_analysis(y_test) # get an analysis of our prediction

def _prevote(self,row):
    return [int(Y_train[i]) for i in row.tolist()] # transform a row indices
    ↪to its label

def _predict(self,test:tuple):
    test_x_num, test_x_ohe = test
    N =len(test_x_num)
    batch_size = N / self.N_batch # set the batch size
    dataframes_indices = []

    for i in tqdm(range(self.N_batch)): # iterate over each batch
        # free up the GPU RAM
        cp.get_default_memory_pool().free_all_blocks()
        # get the bornes from the batch index
        a,b= round(batch_size*i),round(batch_size*(i+1))
        # get the values from the interval
        num,ohe = test_x_num[a:b],test_x_ohe[a:b]
        # compute the distances between the current test batch and all the
    ↪training set
        temp = self._compute_distance(num,self.x_num) + self.ohe_func(self.
    ↪_compute_distance(ohe,self.x_ohe))
        # get the top indices based on the closets distance
        top_K_indices = cp.argsort(temp, axis=1)[: , :self.K] # TODO check give
    ↪the label now
        dataframes_indices.append(pd.DataFrame(top_K_indices.get()))
        del temp, top_K_indices
        collect() # free up the ram

        sleep(0.1)
    return dataframes_indices

def _to_one_hot_encoding(x):

```

```

        return 0 if x == 0 else 1 # if the difference between the values is not 0
↳ then we set it to 1 since they equally important state

def _compute_distance(self,a,b):
    # compute the distance of all the vector in two matrix

    A = a.to_numpy(dtype='float32')
    B = b.to_numpy(dtype='float32')
    A = cp.asarray(A)
    B = cp.asarray(B)
    A_sq_norms = cp.sum(A ** 2, axis=1).reshape(-1, 1) # Shape (n, 1)
    B_sq_norms = cp.sum(B ** 2, axis=1).reshape(1, -1) # Shape (1, m)

    dot_product = cp.dot(A, B.T) # Shape (n, m)
    euclidean_distances = A_sq_norms + B_sq_norms - 2 * dot_product
    del A_sq_norms,B_sq_norms, dot_product, A,B
    collect()
    return euclidean_distances

def _vote_majority(self,label_vectors):
    # vote the label class based on majority occurrences

    sum_zero, sum_one,n = self._label_count(label_vectors) # BUG
    self.Y_PredProba.append(sum_one/n if sum_one > sum_zero else sum_zero/n)
    return 1 if sum_one > len(label_vectors)-sum_one else 0

```

1.6 Training

```

[35]: def print_accuracy(accuracy):
        return print(f'Accuracy: {accuracy:.4f}')

def my_model_validator(best_model : DecisionTreeClassifier,d,s,i,c): # function
↳ to train and get the best model based on accuracy between two model

    model = DecisionTreeClassifier(d,i,s,c)
    model.fit(x_train, y_train)
    model.train(x_val,y_val)

    if best_model is None:
        return model

    if model > best_model:
        best_model = model

    return best_model

```

```

def scikit_learn_val(best_model,d,s,i,c):# function to train and get the best
↳model based on accuracy between two model
    model = DTC(criterion='gini' if c =='gini_index' else
↳'entropy',max_depth=d,min_impurity_decrease=i,min_samples_split=s)
    model.fit(x_train, y_train)
    y_pred = model.predict(x_val)
    accuracy = accuracy_score(y_val, y_pred)
    if best_model is None:
        return model,accuracy
    b_model,b_accuracy = best_model
    if accuracy > b_accuracy:
        b_model = model
        b_accuracy = accuracy

    return b_model,b_accuracy

def train()->tuple[DecisionTreeClassifier,DTC]: # return the best model based
↳on accuracy of scikit learn and my own by different combinaison of parameter
    max_depth= [6,10,12,14]
    min_samples_split= [50,100,300,500]
    min_impurity_decrease=[0.0, 0.001, 0.01]
    criterion:list[ImpurityType]= ['gini_index', 'entropy']

    my_best_model = None
    scikit_best_model = None
    for d in max_depth:
        for s in min_samples_split:
            for i in min_impurity_decrease:
                for c in criterion:
                    my_best_model = my_model_validator(my_best_model,d,s,i,c)
                    scikit_best_model
↳=scikit_learn_val(scikit_best_model,d,s,i,c)

    return my_best_model,scikit_best_model

```

```
[36]: collect()
```

```
[36]: 0
```

Due to a constraint of time and resource I was not able to train the model for all situation, so below are the best model for the highly correlated data removed without PCA

```

[37]: #my_best_model,scikit_best_model= train()
      #scikit_best_model,b_accuracy = scikit_best_model
      #DTC(max_depth=14, min_samples_split=50)
      #Accuracy: 0.9478

```

```
#DecisionTreeClassifier(Max_Depth=14,Min_Inf_Gain=0.  
↪0,Min_Sample=50,Impurity=gini_index)  
#Accuracy: 0.9457
```

1.7 Testing

1.7.1 Model Testing

Lets try another model of My DecisionTreeClassifier

```
[68]: my_DTC = DecisionTreeClassifier(14,0.001,100,'gini_index',)  
my_DTC.fit(X_train,Y_train)  
my_DTC.train()  
my_DTC.predict(X_test,Y_test)  
print_accuracy(my_DTC.accuracy)
```

Accuracy: 0.7981

ScikitLearn DecisionTree

```
[74]: scikit_DTC = DTC(max_depth=14,min_samples_split=50)  
scikit_DTC.fit(X_train,Y_train)  
Y_pred = scikit_DTC.predict(X_test)  
accuracy = accuracy_score(Y_test, Y_pred)  
print_accuracy(accuracy)
```

Accuracy: 0.7331

ScikitLearn KNN

```
[73]: knn = KNeighborsClassifier(n_neighbors=419, algorithm='brute')  
knn.fit(X_train, Y_train)  
y_pred = knn.predict(X_test)  
accuracy = accuracy_score(Y_test, y_pred)  
print_accuracy(accuracy)
```

Accuracy: 0.8181

My KNN from scratch

```
[72]: my_KNN = KNNClassifier(features_to_ohe,419,N_batch=102)  
my_KNN.fit(X_train,Y_train)  
my_KNN.predict(X_test,Y_test)  
print_accuracy(my_KNN.accuracy)
```

100%| | 102/102 [02:04<00:00, 1.22s/it]

Accuracy: 0.8219


1.7.2 Model Selection


In this scenario, retaining the correlated features, rather than reducing the data to a 30-dimensional PCA space, resulted in my model generally outperforming the Scikit-learn implementation. For the Decision Tree algorithm, after training on the `X_train` and `X_val` datasets and selecting the optimal model for the `X_test` dataset, my model showed slightly better performance. Similarly, for the KNN algorithm, both models performed similarly, with mine achieving a slight edge. Moving forward, we will compare my models across additional metrics to identify the most effective one.


```
[69]: def show_metrics(model:BinaryClassifier):  
    print(f"Accuracy: {model.accuracy*100:.5f} %")  
    print(f'Precision: {model.precision*100:.5f} %')  
    print(f'Recall: {model.recall*100:.5f} %')  
    print(f'F-score: {model.f_score*100:.5f} %')  
    print(f'')  
    model.plot_confusion_matrix()  
    print(f'')  
    model.plot_roc_curve()  
    print(f'')  
    model.plot_precision_recall_curve()  
    print(f'')
```

```
[71]: show_metrics(my_DTC)
```

```
Accuracy: 79.81465 %  
Precision: 92.12865 %  
Recall: 60.22973 %  
F-score: 72.83989 %
```










```
[75]: show_metrics(my_KNN)
```

```
Accuracy: 82.18676 %  
Precision: 96.80610 %  
Recall: 62.42162 %  
F-score: 75.90128 %
```





1.8 Conclusion

```
[88]: # Data
scenarios = ['Removed w/o PCA', 'Removed w/ PCA', 'Kept w/o PCA', 'Kept w/ PCA']

# Accuracies for each scenario
removed_no_pca = [0.7962, 0.8401, 0.8066, 0.7979]
removed_pca = [0.7371, 0.7320, 0.7241, 0.7239]
kept_no_pca = [0.7981, 0.7374, 0.8181, 0.8219]
kept_pca = [0.7905, 0.6500, 0.7674, 0.7675]

# X-axis labels
x_labels = ['MyDTC', 'SlDTC', 'SlKNN', 'My DTC']

data = [removed_no_pca, removed_pca, kept_no_pca, kept_pca]

fig, axs = plt.subplots(2, 2, figsize=(16, 12)) # Increase figure size for
↳ more space
```

```

fig.suptitle('Accuracy Across Different Scenarios', fontsize=18)

titles = ['Removed Correlated Feature without PCA',
          'Removed Correlated Feature with PCA',
          'Correlated Features Kept without PCA',
          'Correlated Features Kept with PCA']

for i, ax in enumerate(axes.flat):
    ax.plot(x_labels, data[i], marker='o', color='b') # Use custom x_labels
    ↪for x-axis
    ax.set_title(titles[i])
    ax.set_xlabel('Model Type')
    ax.set_ylabel('Accuracy')
    ax.set_ylim([0.6, 0.85]) # Adjust the limits for y-axis for better
    ↪comparison
    ax.grid(True)

    ax.set_xticklabels(x_labels, rotation=30, ha='right')

    for j, v in enumerate(data[i]):
        ax.text(j, v + 0.005, f"{v:.4f}", ha='center', fontweight='bold')


plt.show()

```

<ipython-input-88-b3a3936fa8f9>:39: UserWarning: FixedFormatter should only be used together with FixedLocator

```
ax.set_xticklabels(x_labels, rotation=30, ha='right')
```

Accuracy Across Different Scenarios



In the graphics above, we observe that when the highly correlated features were removed, the Scikit-learn model demonstrated better predictions, suggesting it performed faster and more efficiently with fewer features. However, overall, the mean accuracy across all scenarios was higher for my model compared to Scikit-learn's, with mine achieving 78% accuracy versus Scikit's 76%. While we could choose to remove the highly correlated features and retain the Scikit model for comparison, my model remains slightly better overall—and I prefer it, as it allowed for easier computation of additional metrics.

We ultimately concluded that the KNN model was the most effective due to its superior precision and F-score. For further details, please refer to the report!